



APACHE
kafka®

A distributed streaming platform

Streaming Platform?

pub/sub similar to
message queues



Store streams of records
in a **fault-tolerant** &
durable way

Process streams of
records as they **occur**

Typical Use Cases?

Messaging

Website Activity Tracking

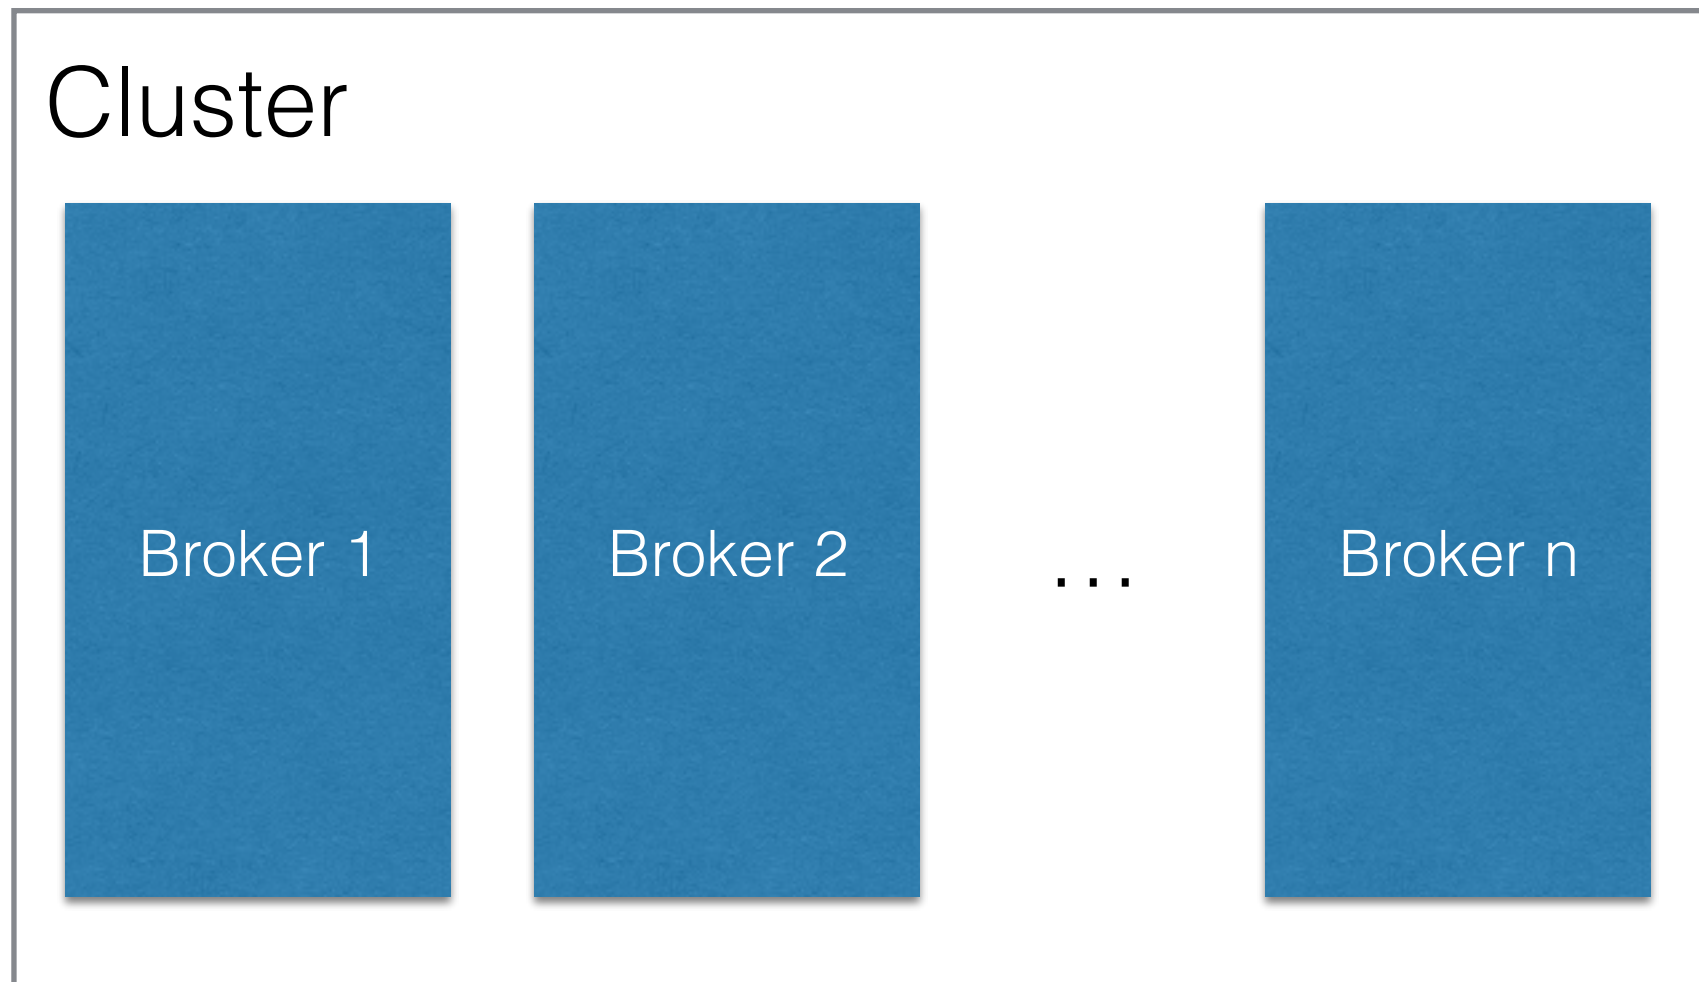
Log Aggregation

Stream Processing

Event Sourcing


Core Concepts

Cluster & Broker

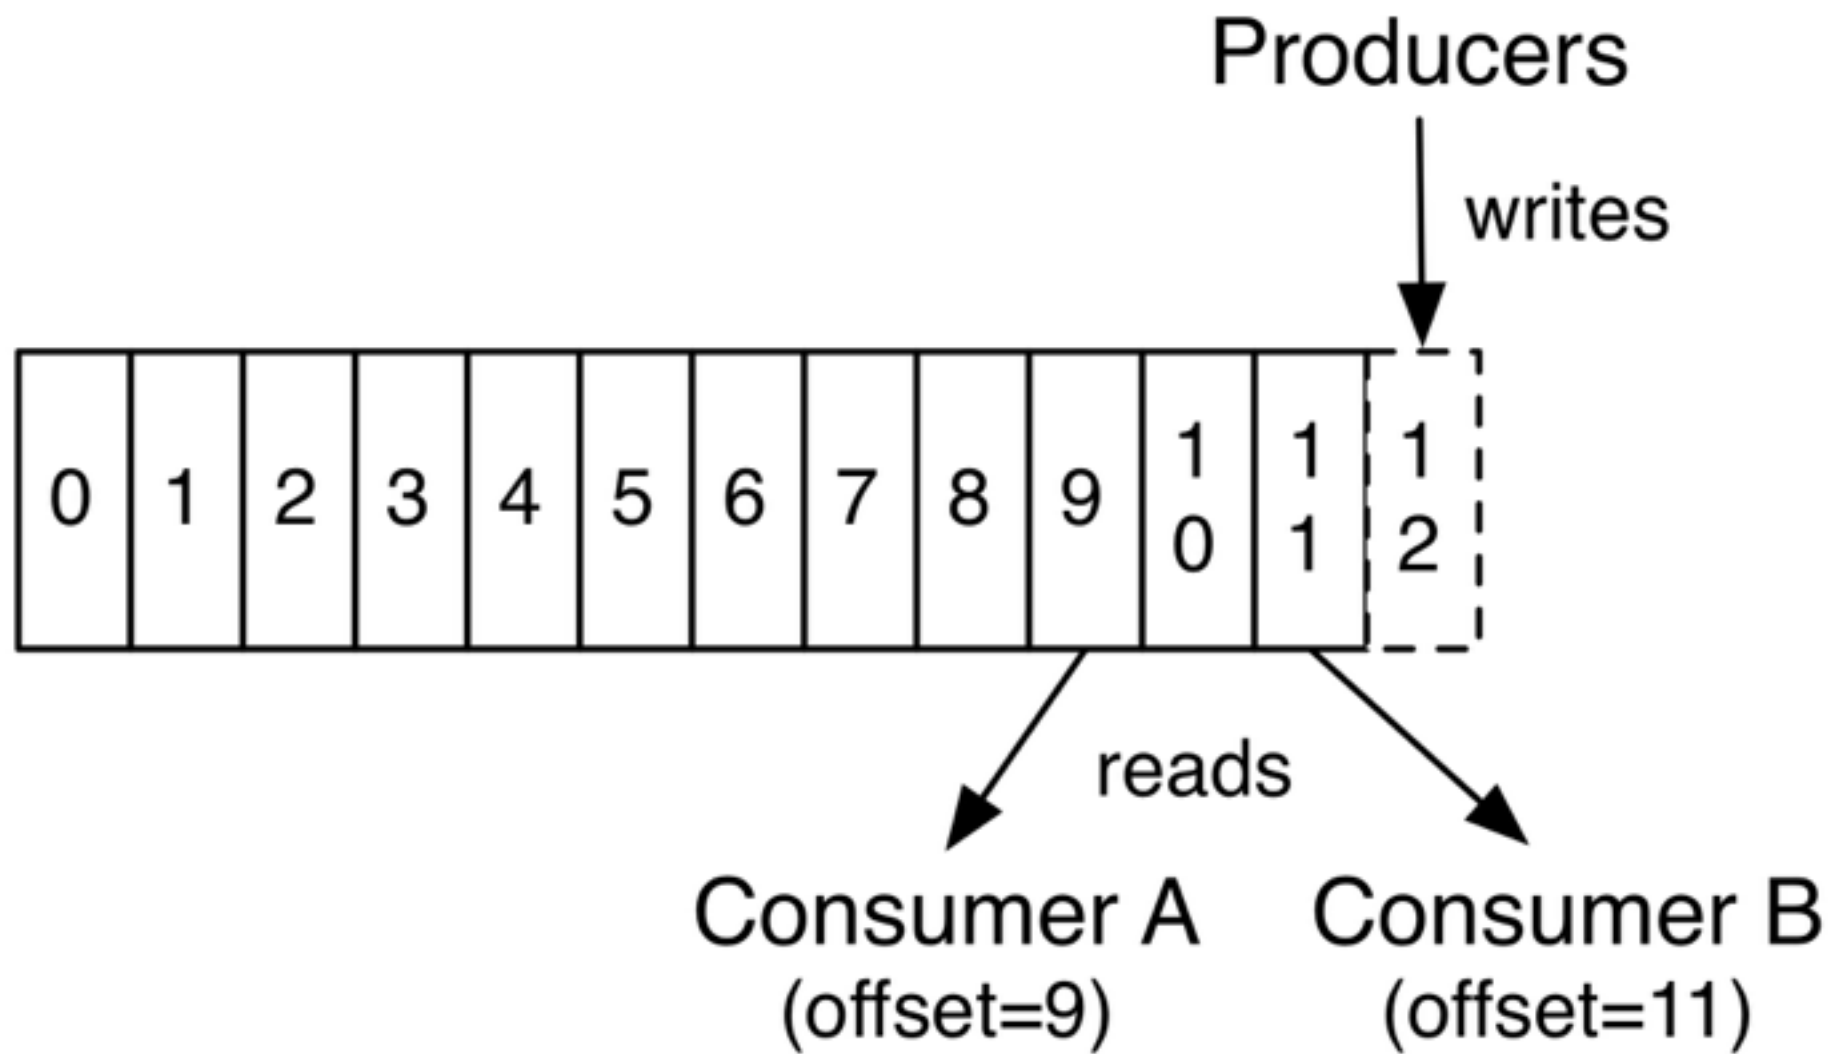


Topic

Anatomy of a Topic



Partition



Record

key;value;timestamp

Retention

- By time (e.g. max 2 days)
- By size (log may not grow beyond X GB)
- Infinite (totally fine)

Log compaction

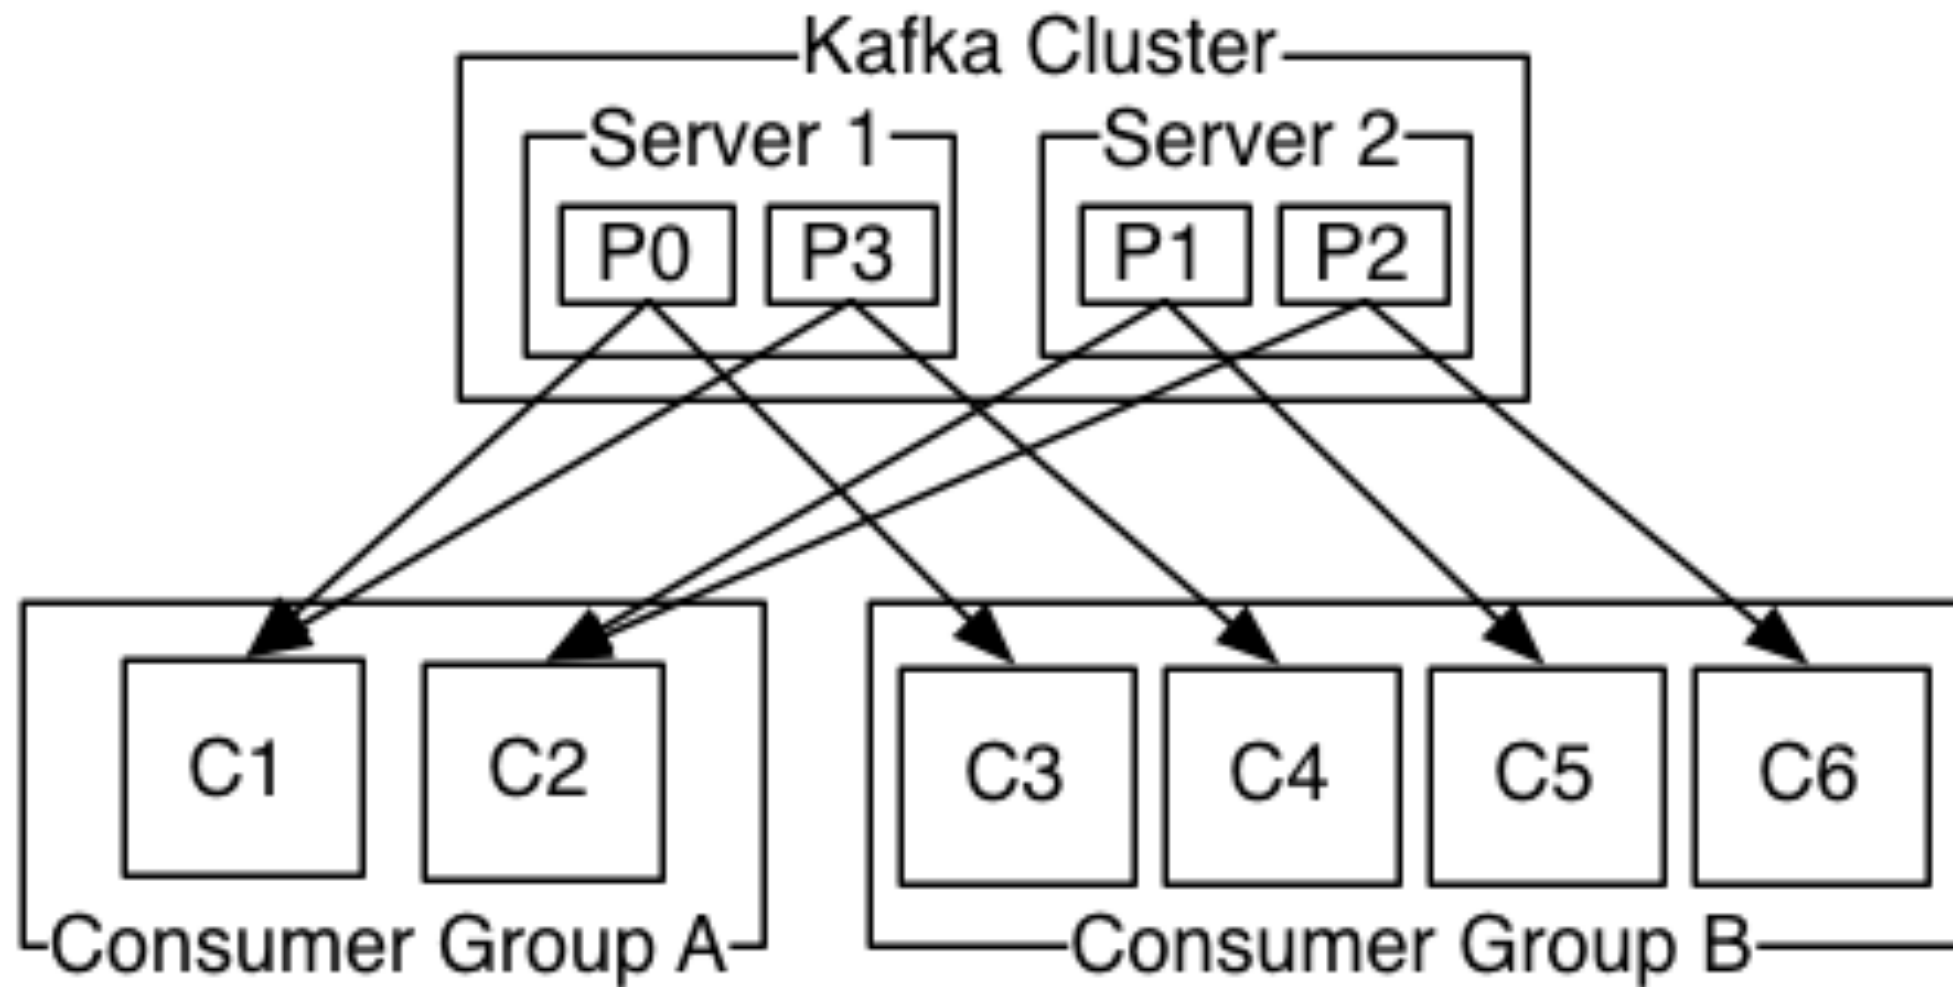
```
1 123 => bill@microsoft.com  
2      .  
3      .  
4      .  
5 123 => bill@gatesfoundation.org  
6      .  
7      .  
8      .  
9 123 => bill@gmail.com
```

Quotas

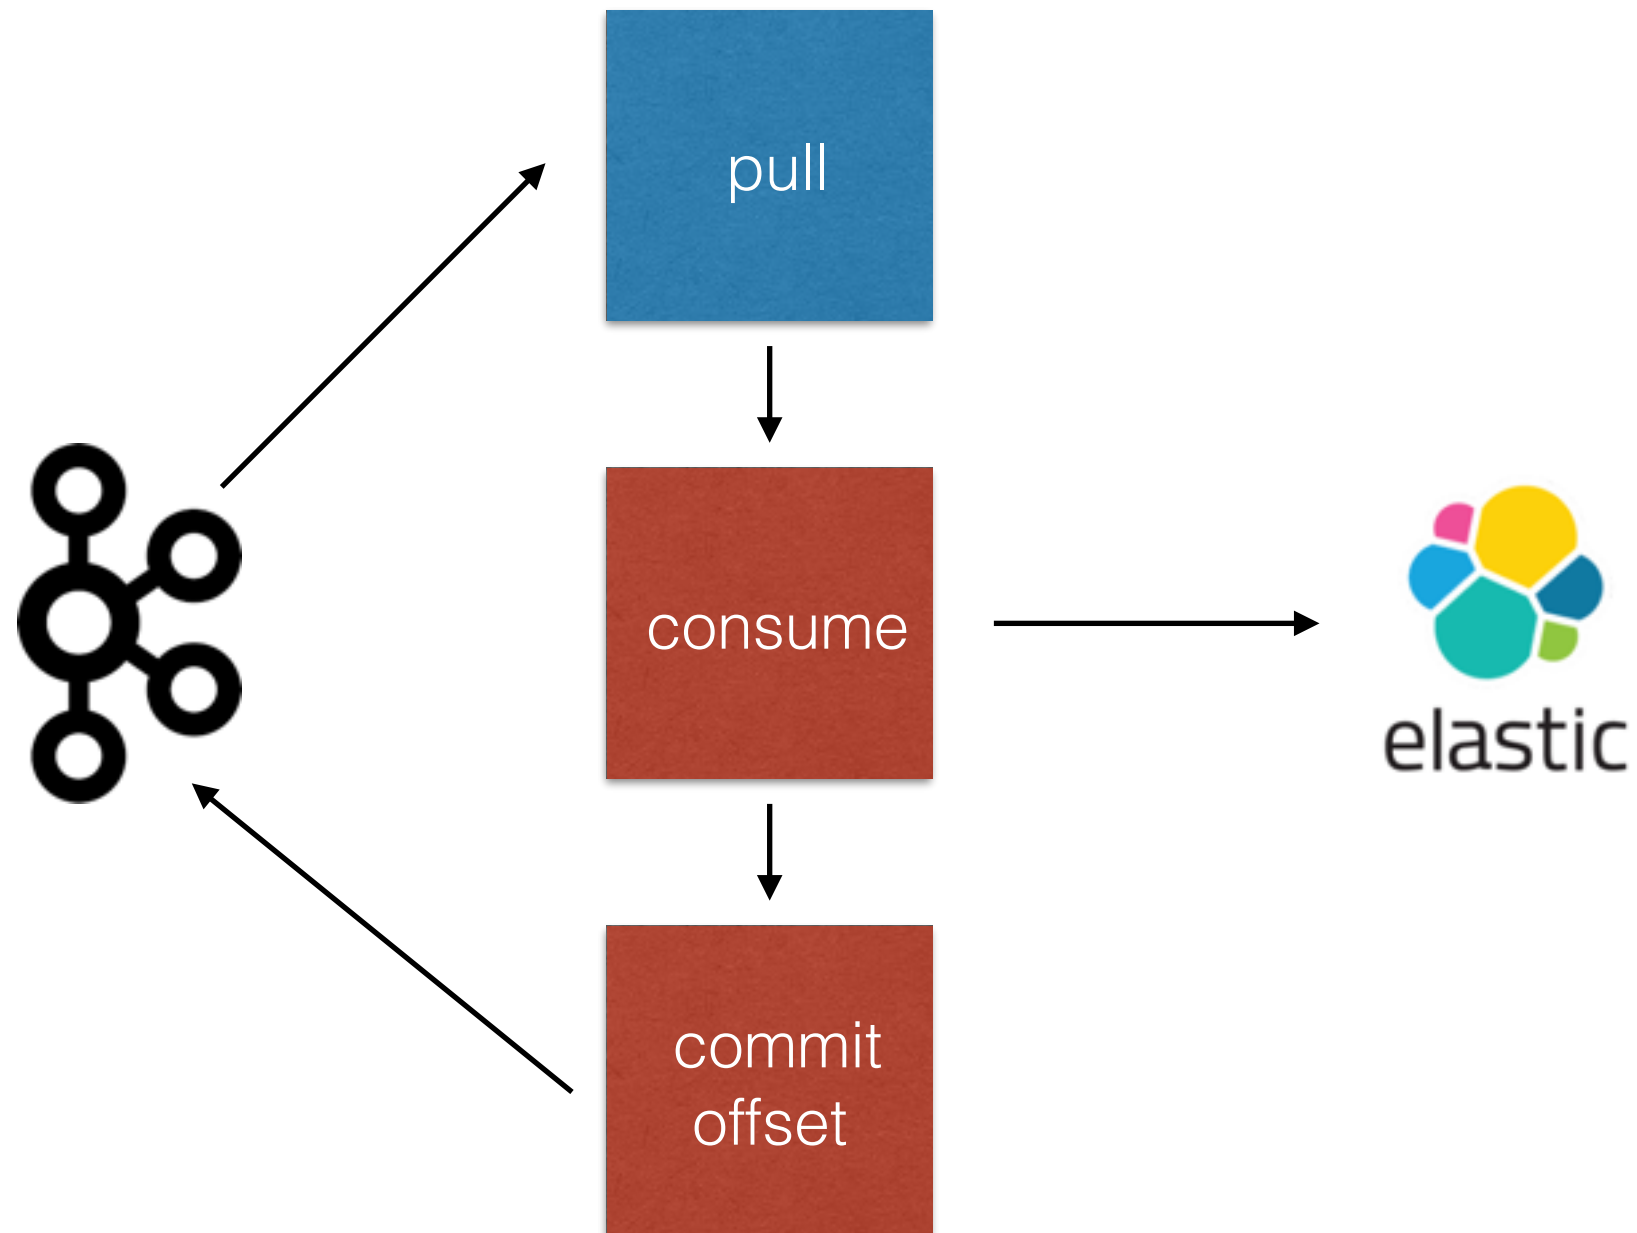
- Network bandwidth
- Request rate

Producer

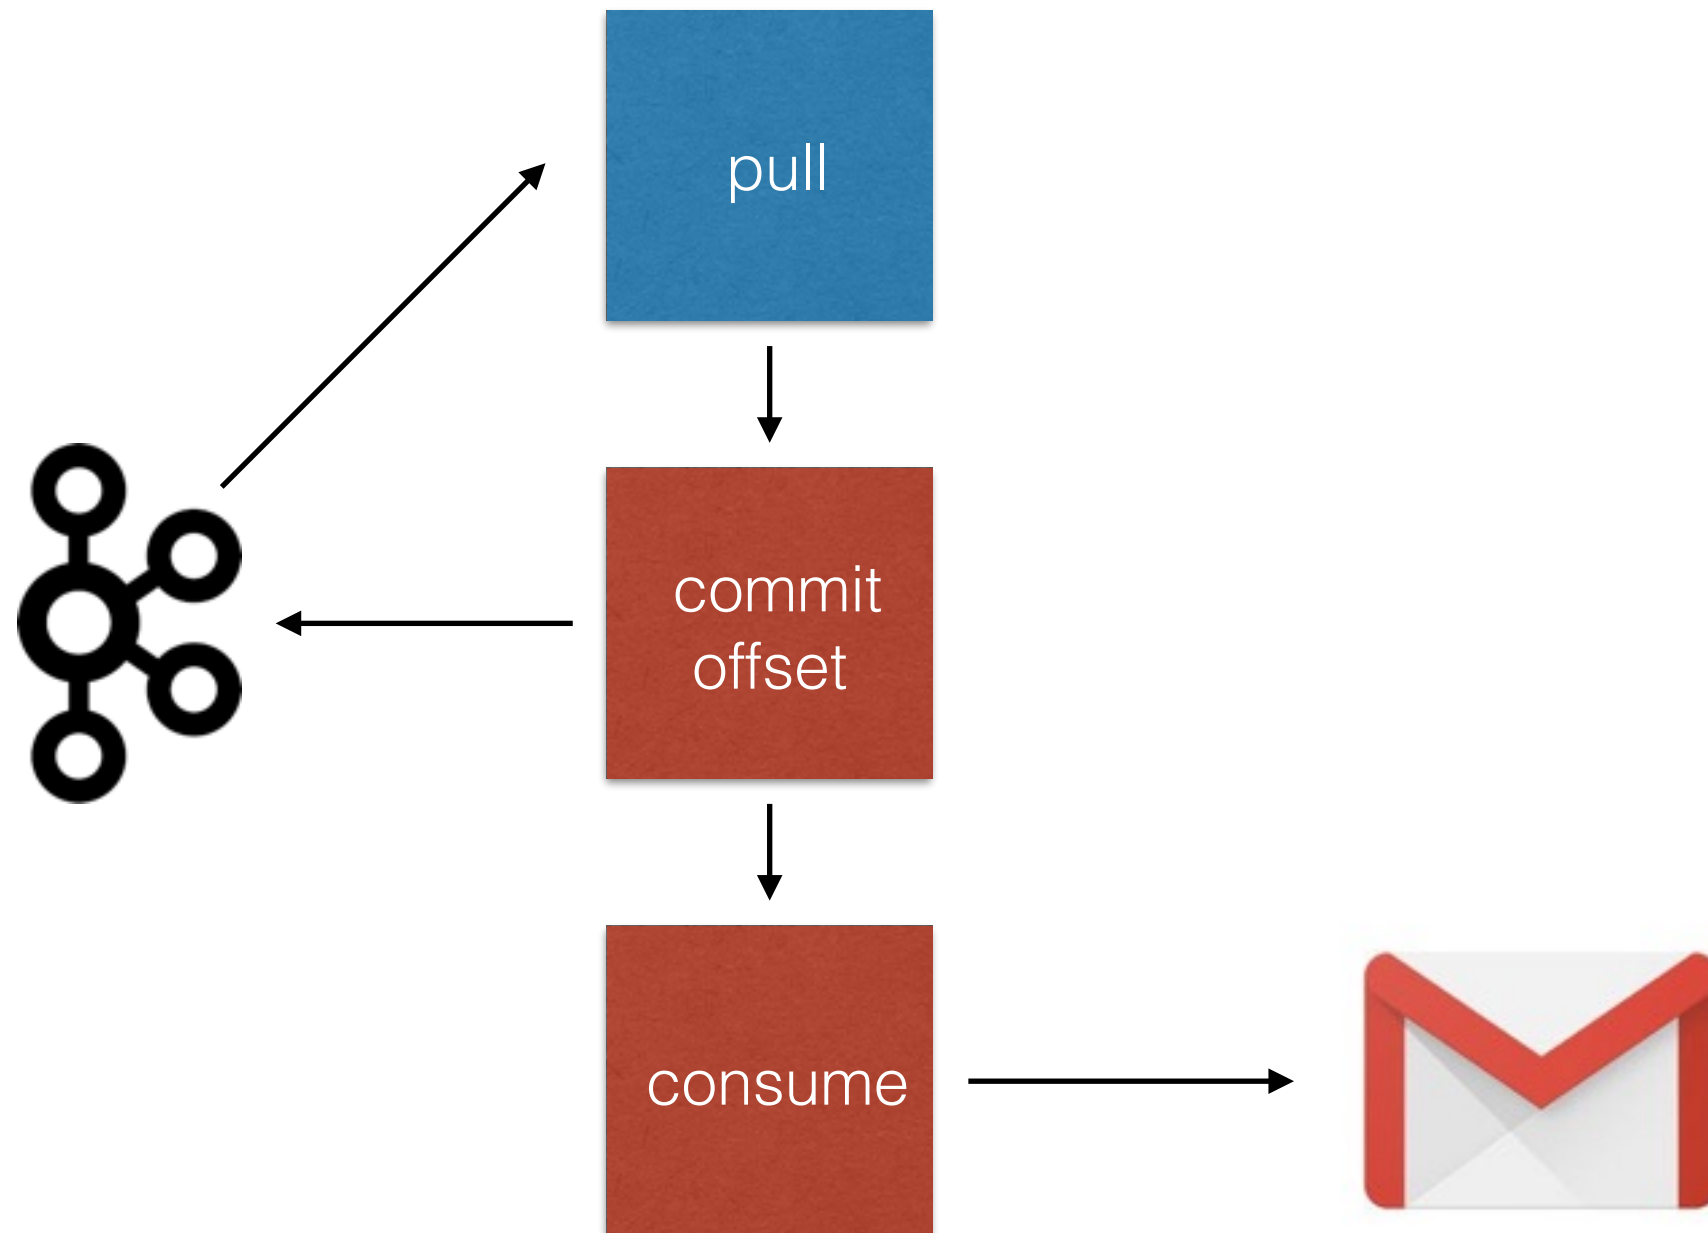
Consumer



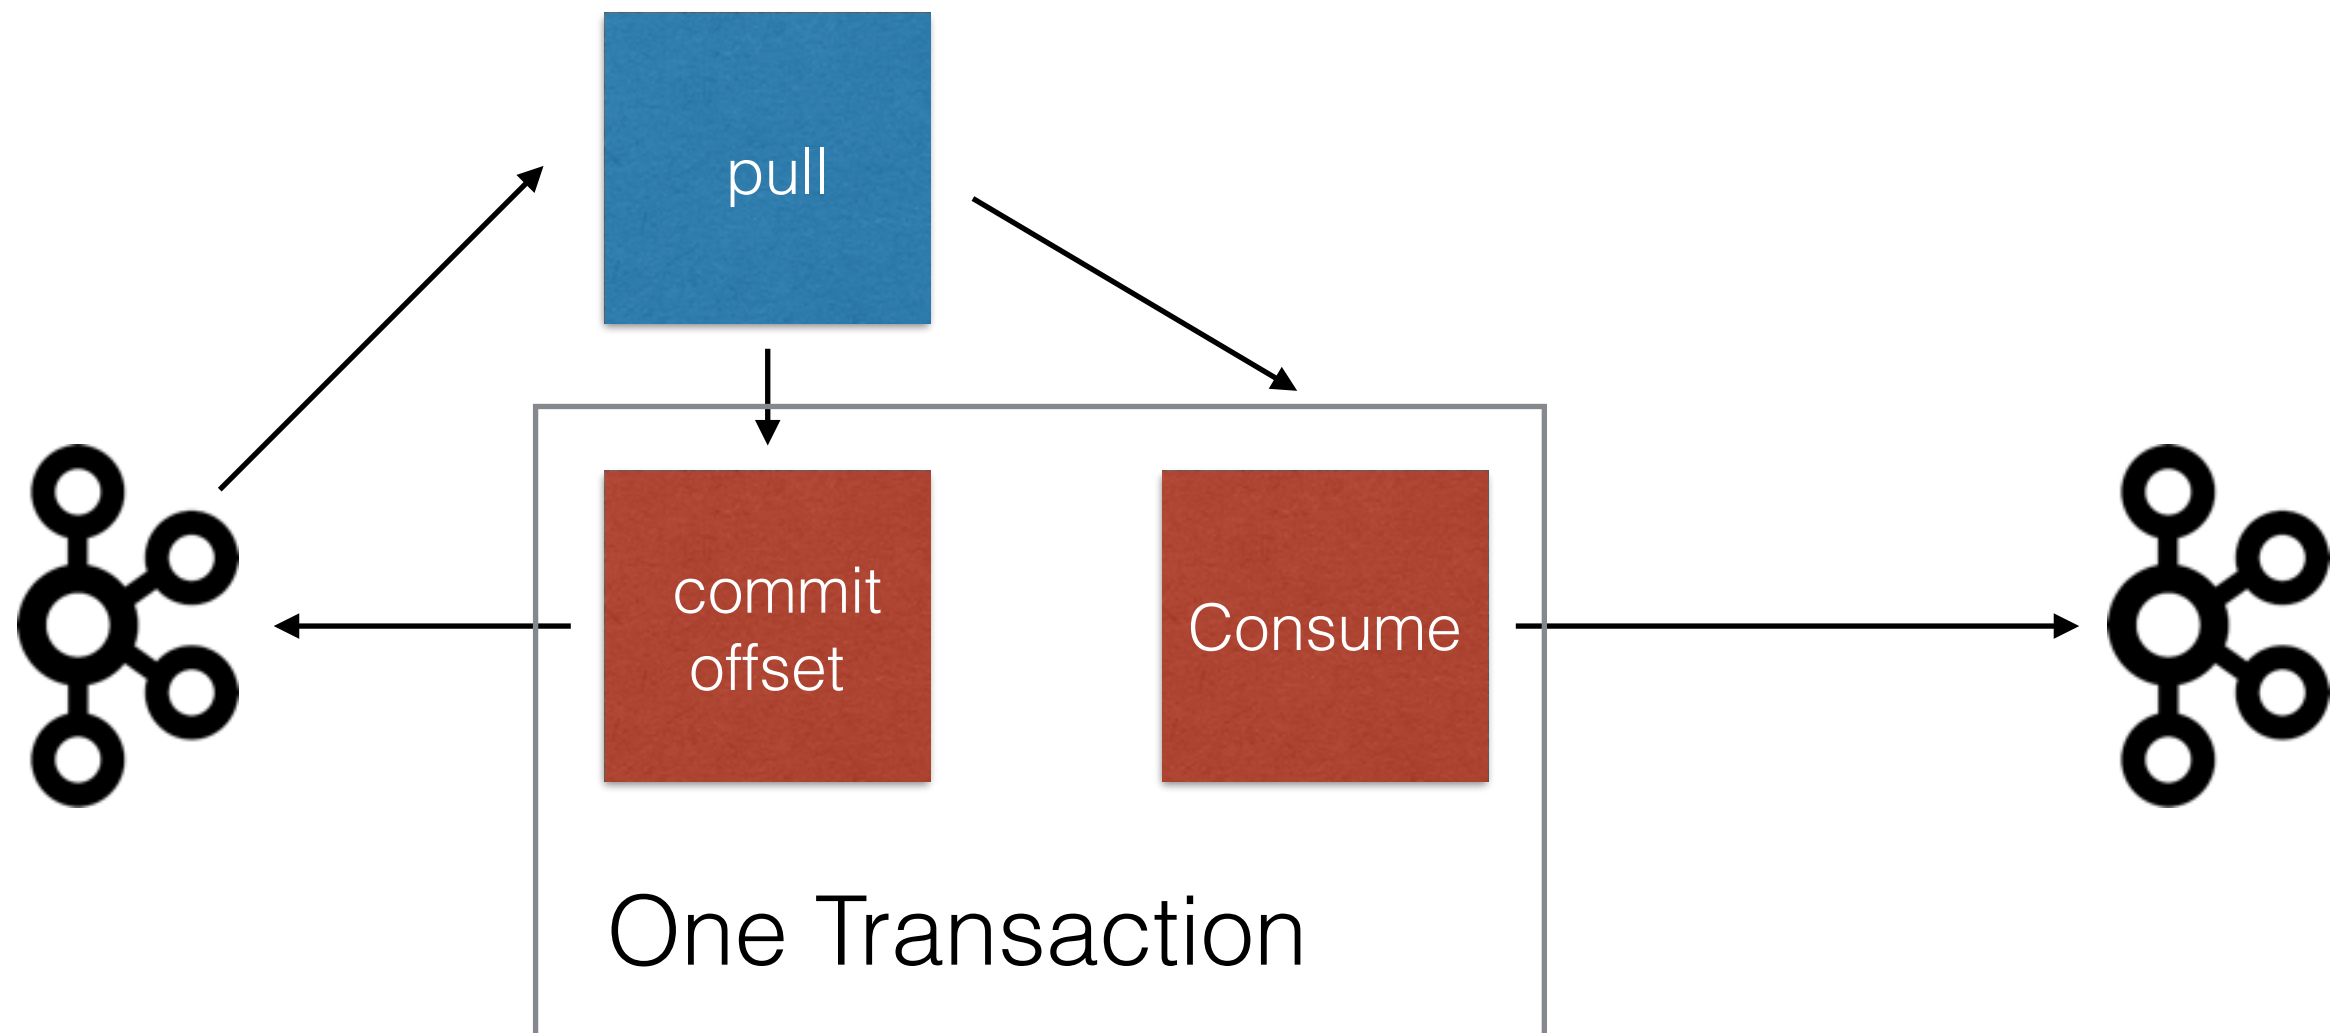
At least once



At most once



Exactly once



Let's code!

[https://github.com/
Sebruck/kafka-talk](https://github.com/Sebruck/kafka-talk)